

S3:

What Is S3?



Object Storage

S3 provides secure, durable, highly scalable object storage.



Scalable

S3 allows you to store and retrieve any amount of data from anywhere on the web at a very low cost.



Simple

Amazon S3 is easy to use, with a simple web service interface.

S3 IS OBJECT-BASED STORAGE

Manages data as objects rather than in file systems or data blocks.

- ✓ **Upload any file type** you can think of to **S3**.
- ✓ **Examples** include photos, videos, code, documents, and text files.
- ✓ **Cannot** be used to run an operating system or database.

S3 Basics



Unlimited Storage

The total volume of data and the number of objects you can store is unlimited.



Objects up to 5 TB in Size

S3 objects can range in size from a minimum of 0 bytes to a maximum of 5 terabytes.



S3 Buckets

Store files in buckets (similar to folders).

Working with S3 Buckets

1 Universal Namespace


All AWS accounts share the S3 namespace.
Each S3 bucket name is globally unique.

2 Example S3 URLs

<https://bucket-name.s3.Region.amazonaws.com/key-name>
<https://acloudguru.s3.us-east-1.amazonaws.com/Ralphie.jpg>

3 Uploading Files

When you upload a file to an S3 bucket, you will receive an HTTP 200 code if the upload was successful.



Key

The name of the object
(e.g., `Ralphie.jpg`)



Version ID

Important for storing multiple
versions of the same object



Value

The data itself, which is made
up of a sequence of bytes



Metadata

Data about the data you are
storing (e.g., `content-type`,
`last-modified`, etc.)

S3 is a safe place to store your files.

The data is spread across multiple devices and
facilities to **ensure availability** and **durability**.



Highly Available and Highly Durable



Built for Availability

Built for 99.95% – 99.99% **service availability**, depending on the S3 tier.



Designed for Durability

Designed for 99.999999999% (9 decimal
places) durability for **data stored** in S3.

S3 Standard

1 High Availability and Durability

Data is stored redundantly across multiple devices in multiple facilities (≥ 3 AZs):

- 99.99% availability
- 99.999999999% durability (11 9's)

2 Designed for Frequent Access

Perfect for frequently accessed data.

3 Suitable for Most Workloads

- The default storage class.
- Use cases include websites, content distribution, mobile and gaming applications, and big data analytics.



Tiered Storage

S3 offers a range of storage classes designed for different use cases.

Lifecycle Management

Define rules to automatically transition objects to a cheaper storage tier or delete objects that are no longer required after a set period of time.

Versioning

With versioning, all versions of an object are stored and can be retrieved, including deleted objects.

Securing Your Data

1 Server-Side Encryption

You can set default encryption on a bucket to encrypt all new objects when they are stored in the bucket.


2 Access Control Lists (ACLs)

Define which AWS accounts or groups are granted access and the type of access. You can attach S3 ACLs to individual objects within a bucket.

3 Bucket Policies

S3 bucket policies specify what actions are allowed or denied (e.g., allow user Alice to PUT but not DELETE objects in the bucket).





Strong Read-After-Write Consistency

- **After a successful write** of a new object (PUT) or an overwrite of an existing object, any subsequent read request immediately receives the latest version of the object.
- **Strong consistency** for list operations, so after a write, you can immediately perform a listing of the objects in a bucket with all changes reflected.

What to Know for the Exam



Object Based

Object-based storage allows you to upload files.



Not OS or DB Storage

Not suitable to install an operating system or run a database on.



Files up to 5 TB

Files can be from 0 bytes to 5 TB.



Unlimited Storage

The total volume of data and the number of objects you can store is unlimited.

💡 IMPORTANT EXAM TIP

Files Stored in Buckets

S3 is a **universal namespace**.

<https://bucket-name.s3.Region.amazonaws.com/key-name>
<https://acloudguru.s3.us-east-1.amazonaws.com/Ralphie.jpeg>

- ✓ Successful CLI or API uploads will generate an **HTTP 200 status code**.

4 S3 Object Tips

- ✓ **Key**
The object name (e.g., `Ralphie.jpg`)
- ✓ **Value**
The data itself, which is made up of a sequence of bytes
- ✓ **Version ID**
Allows you to store multiple versions of the same object
- ✓ **Metadata**
Data about the data you are storing (e.g., `content-type`, `last-modified`, etc.)